

Task no: 1

Code:

```
# =====  
# Text Classification with BERT (Complete & Correct)  
# =====  
# Install libraries  
!pip install transformers datasets accelerate -q  
# 1. Imports  
import torch  
import numpy as np  
import random  
from torch.utils.data import DataLoader  
from torch.optim import AdamW  
from transformers import (  
    AutoTokenizer,  
    AutoModelForSequenceClassification,  
    get_scheduler,  
    DataCollatorWithPadding  
)  
from datasets import load_dataset  
from sklearn.metrics import accuracy_score  
# Set seed for reproducibility  
seed = 42  
torch.manual_seed(seed)
```

```

np.random.seed(seed)

random.seed(seed)

if torch.cuda.is_available():
    torch.cuda.manual_seed_all(seed)

# 2. Load Dataset (AG News)

dataset = load_dataset("ag_news")

train_dataset = dataset["train"].select(range(2000)) # Smaller subset for quick
training

test_dataset = dataset["test"].select(range(500))    # Smaller test set

# 3. Tokenizer

tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")

def tokenize_fn(batch):
    return tokenizer(batch["text"], truncation=True, max_length=256)

# Tokenize datasets

train_dataset = train_dataset.map(tokenize_fn, batched=True)

test_dataset = test_dataset.map(tokenize_fn, batched=True)

# Rename column

train_dataset = train_dataset.rename_column("label", "labels")

test_dataset = test_dataset.rename_column("label", "labels")

# Set format

train_dataset.set_format(type="torch", columns=["input_ids", "attention_mask",
"labels"])

test_dataset.set_format(type="torch", columns=["input_ids", "attention_mask",
"labels"])

# 4. DataLoaders with dynamic padding

data_collator = DataCollatorWithPadding(tokenizer=tokenizer)

```

```
train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True,  
collate_fn=data_collator)
```

```
test_loader = DataLoader(test_dataset, batch_size=16,  
collate_fn=data_collator)
```

5. Model

```
model = AutoModelForSequenceClassification.from_pretrained(  
    "bert-base-uncased",  
    num_labels=4  
)
```

6. Optimizer

```
optimizer = AdamW(model.parameters(), lr=2e-5) # Slightly lower learning  
rate
```

7. Training Setup

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")  
print(f"Using device: {device}")  
model.to(device)  
num_epochs = 1  
num_training_steps = len(train_loader) * num_epochs  
lr_scheduler = get_scheduler(  
    "linear",  
    optimizer=optimizer,  
    num_warmup_steps=0,  
    num_training_steps=num_training_steps  
)
```

8. Training Loop

```
model.train()  
print("Starting training...")
```

```

for epoch in range(num_epochs):
    total_loss = 0
    for step, batch in enumerate(train_loader):
        batch = {k: v.to(device) for k, v in batch.items()}
        outputs = model(**batch)
        loss = outputs.loss
        loss.backward()
        optimizer.step()
        lr_scheduler.step()
        optimizer.zero_grad()
        total_loss += loss.item()
        if (step + 1) % 50 == 0:
            avg_loss = total_loss / 50
            print(f'Step {step + 1}/{len(train_loader)} - Loss: {avg_loss:.4f}')
            total_loss = 0
print("✅ Training Finished!")

```

9. Evaluation

```

model.eval()
all_preds = []
all_labels = []
print("Running evaluation...")
with torch.no_grad():
    for batch in test_loader:
        batch = {k: v.to(device) for k, v in batch.items()}
        outputs = model(**batch)
        preds = torch.argmax(outputs.logits, dim=-1)

```

```
all_preds.extend(preds.cpu().numpy())  
all_labels.extend(batch["labels"].cpu().numpy())  
acc = accuracy_score(all_labels, all_preds)  
print(f"✅ Test Accuracy: {acc:.4f}")  
print(f"✅ Test Samples: {len(all_labels)}")
```